

ML for Memory Allocation

Jan 22 2026

Problem setting

- Memory allocators (e.g., glibc, tcmalloc, jemalloc, Hoard)
 - Performance vs memory consumption
 - Optimizations: per-thread buffers, arenas
 - Object-size classes to avoid bloat
- Allocation lifetime also important
 - Long lived objects cause hugepage fragmentation
- C++ vs managed runtimes (garbage collection)

Challenge

- How to predict allocation lifetime?
- Naïve approach:
 - Profile an application – record lifetimes at each allocation site
 - Drawback: not generalizable (revisions, variations across executions)
- Online approach:
 - Overhead
 - Unseen contexts

LLAMA contributions

- ML to predict lifetime using context
- Challenges:
 - Training? How to generalize?
 - Prediction overheads
 - Mispredictions
- Follow-up work: TelaMalloc allocation for ML accelerators

Discussion

1. When is ML a “viable” tool for memory allocator? What signals make the problem “learnable” here? When it make it a bad fit?
2. The design uses many constants (e.g., K, M, lifetime classes, region size) . Are there any benefits to tuning them? Is yes, how would the values affect performance and fragmentation?

Title: [Lecture 4] Discussion