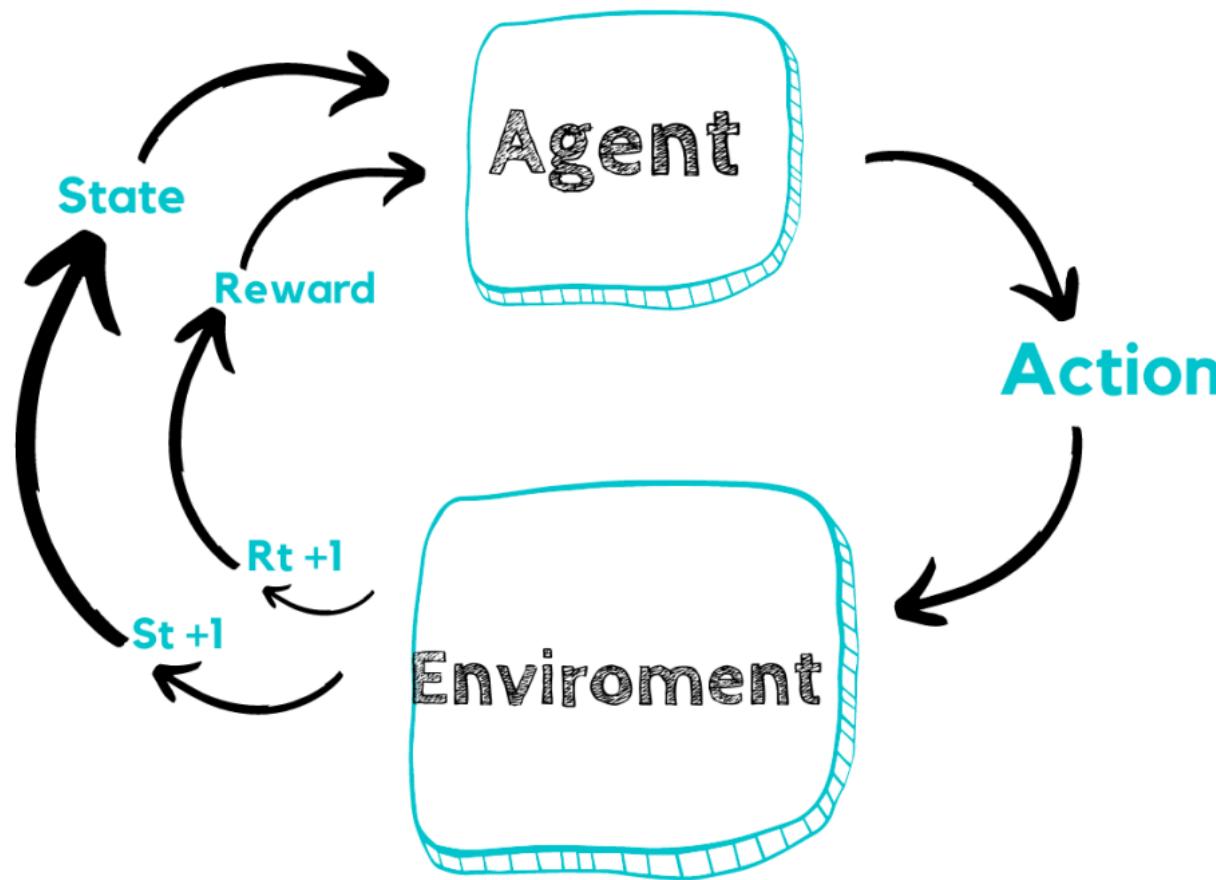


# ArtMem: RL for memory tiering

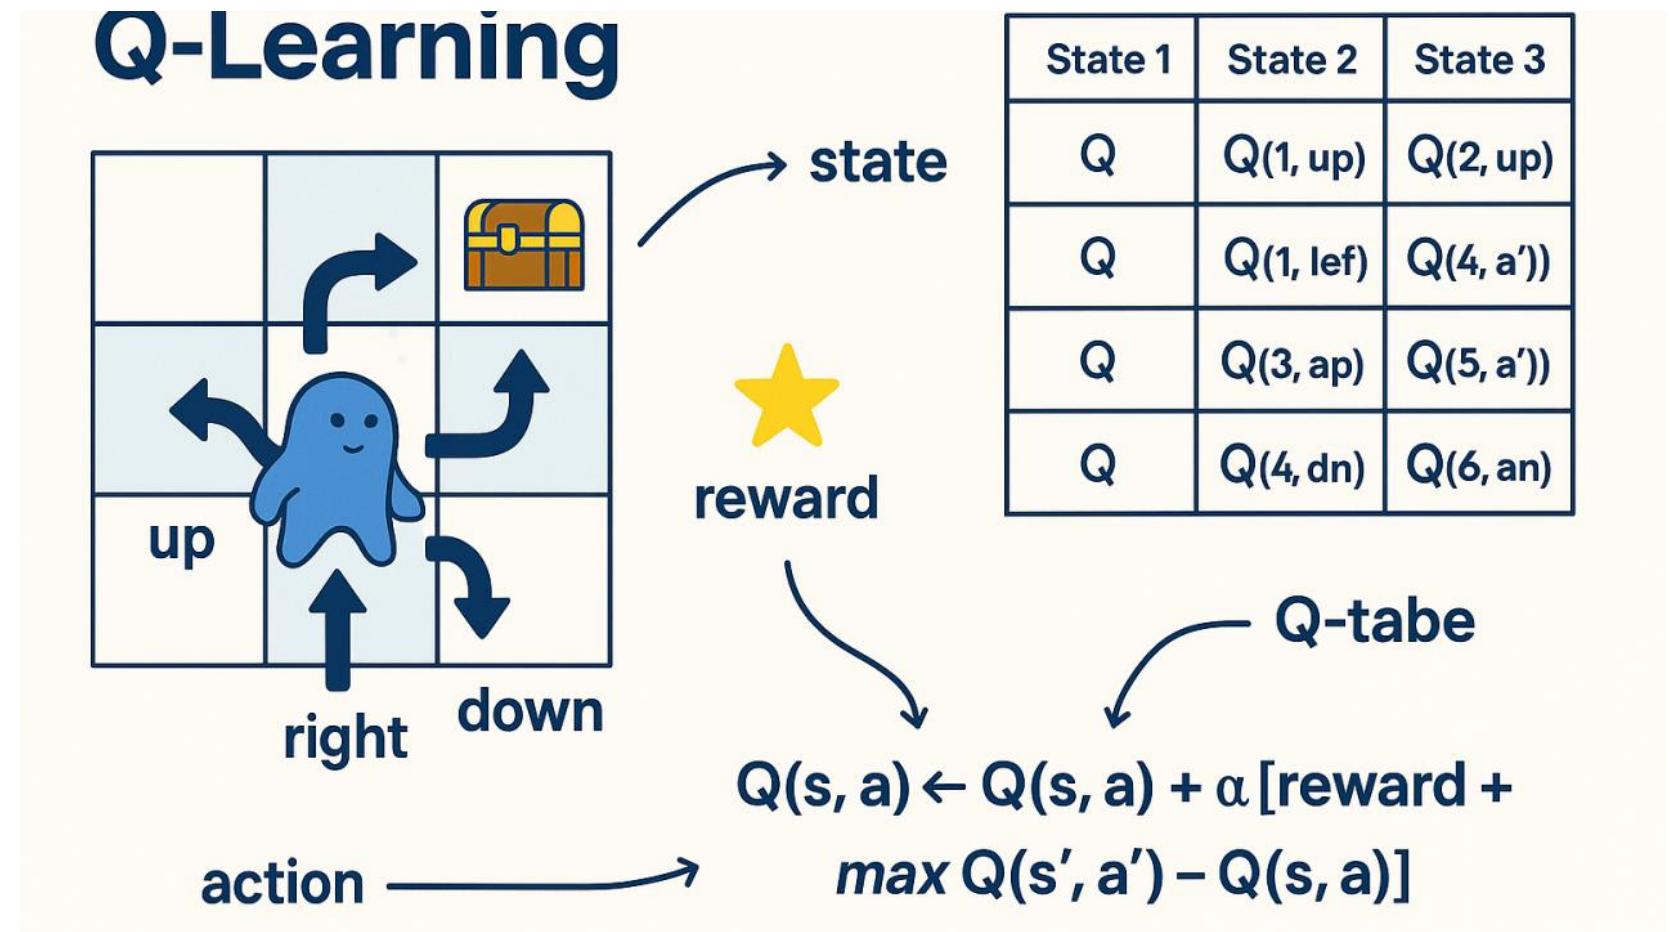
Jan 29

# Reinforcement Learning



# Q-learning

- Value-based
- Model-free
- Off-policy



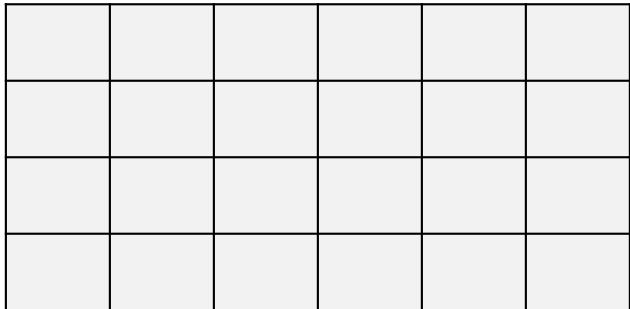
# Memory tiering

- Demand for memory increasing
- Memory ~35% of total rack cost
- Problem: DRAM scaling has stalled; limited pins on chip
- Solution: use memory alternatives
  - E.g. Non-volatile memory (NVM) or CXL
  - Challenge: New tiers are slower

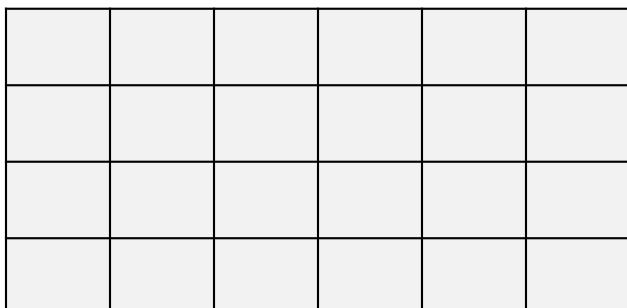
# Tiering systems

- Goal: Minimize accesses to slow tiers; place “hot” data in fast tier and “cold” data in slow tier
- Which pages to place in fast tier?
  - Access frequency or access recency  
E.g.: Page accessed 100 times? Then hot!
- How to adapt to changes in hotness?
  - Migrate pages; either periodically or certain events

# Memory tiering: clever data placement & migration

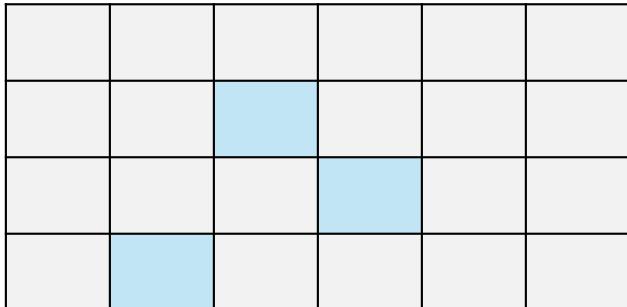


Local DRAM



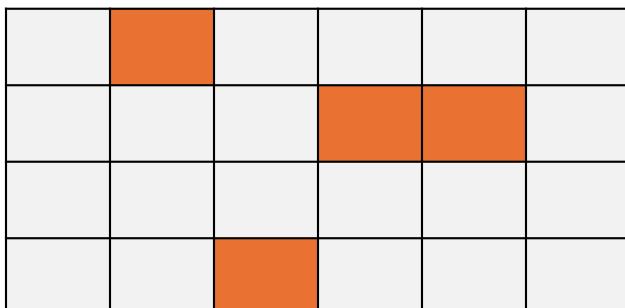
CXL memory

# Memory tiering: clever data placement & migration



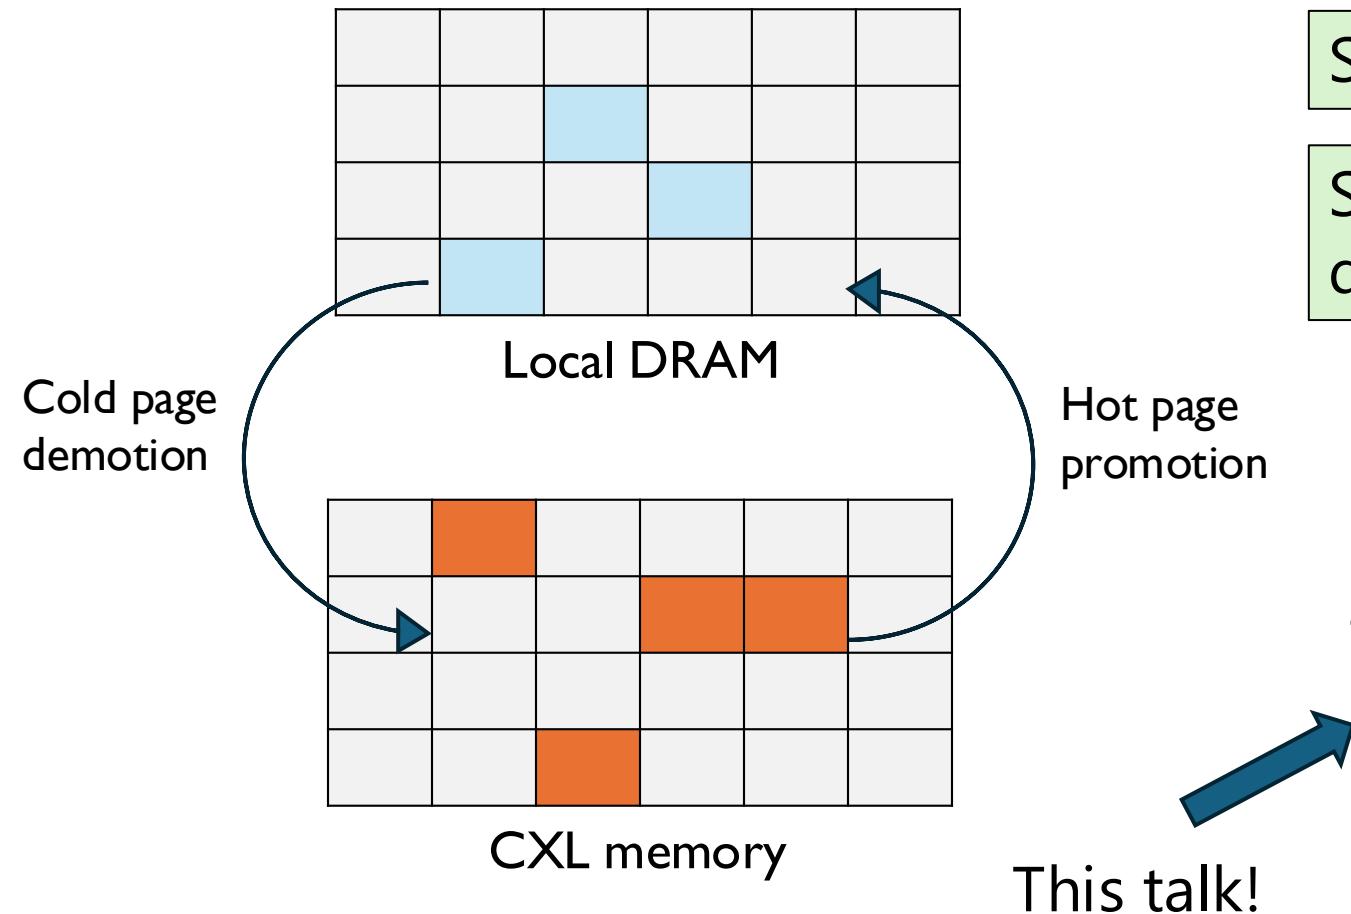
Local DRAM

Step 1: Identify “hot” and “cold” data



CXL memory

# Memory tiering: clever data placement & migration



Step 1: Identify “hot” and “cold” data

Step 2: Promote hot data to local DRAM, demote cold data to CXL memory

Different approaches:

- 1) Hardware based: E.g. Intel Flat MM
- 2) Software based

- Application-transparent (e.g. TPP, HMSDK, Hemem, Memtis)
- Application-aware (e.g. MaPHeA)

# Challenge

- Uses static values
  - How to configure them?

Parameter	Defaults
Hot threshold	8 accesses
Cooling trigger	18 accesses
Sampling rate	1 in 10K
Migration interval	10ms
Migration rate	2 GB/s
....	

# Discussion

1. What are the challenges of RL for systems? When does it work? When might it fail?
2. Another way to build an adaptive tiering system is to use a learned policy directly. The learned policy can predict which pages to migrate and when. Compare this with ArtMem.