

ML deployment challenges

Mar 24 2026

ML OS challenges

- ML inside OS:
 - Low latency constraints (critical path)
 - No access to accelerators
 - Models written in high-level languages
 - Supporting model updates/replacement
 - Framework for data collection for training
 - Correctness and safety guarantees necessary

Discussion

1. Is the feature store in LAKE good enough? Are there any limitations? How could improve it?
2. If you were developing an ML policy inside the OS, what guarantees would you like to provide? How would you do it?